

Bradley_Holt_D212_Task2

January 2, 2022

1 Part I: Research Question

1.1 A1: Organization Question

What customer characteristics or combination of customer characteristics from the available churn_clean data best predict customer retention?

1.2 A2: Goal of Data Analysis

One goal of the data analysis is to reduce the number of features of the dataset. One way to accomplish this is by using feature selection techniques to determine what customer characteristics are similar to each other and dropping one or more from the data set. Another method used is feature extraction and is conducted by combining two or more features into one feature. By reducing the total number of features many benefits are achieved including reduced memory and storage usage and model efficiency.

2 Part II: Method Justification

2.1 B1: Principle Component Analysis

Principal component analysis(PCA) is a fast and flexible unsupervised method for dimensionality reduction in data. Instead of removing entire features from a data set, and therefore information, PCA aims to combine features into one feature by learning about the relationships between them. PCA calculates vectors of the variances between features and transforms them so that said vectors are the x and y axis and scales the data so that variances are uniform.

The expected outcome of running PCA on the dataset is an output of n principle components. Each principle component contains a feature that explains the variation of the original combined variables included in each principle component. The number of principle components can then be selected based on the total cumulative explained variation desired based on desired model accuracy.

2.2 B2: Assumption of Principle Component Analysis

One assumption of PCA is that there is a linear relationship between all variables. The reason for this assumption is that a PCA is based on Pearson correlation coefficients, and as such, there needs to be a linear relationship between the variables. (Laerd Statistics, 2018)

3 Part III: Data Preparation

3.1 C1: Continuous Data Variables from Data Set

Below are the 13 continuous variables to be used to answer the proposed question in A1:

- Lat
- Lng
- Population
- Children
- Age
- Income
- Outage_sec_perweek
- Email
- Contacts
- Yearly_equip_failure
- Tenure
- MonthlyCharge
- Bandwidth_GB_Year

```
[ ]: #importing Pandas and Numpy
import pandas as pd
import numpy as np
```

```
[ ]: #Importing churn_clean data set
churn_clean = pd.read_csv('C:/Users/holtb/Data/D212_Data_Mining_II/data/
    ↪churn_clean.csv')
```

```
[ ]: # Dropping unneed variables
churn_clean.drop(['Customer_id', 'CaseOrder', 'Interaction', 'UID'], axis=1,
    ↪inplace=True)
```

```
[ ]: # Dropping non-continuous variables
non_numeric = ['City', 'State', 'County', 'Zip', 'Area', 'TimeZone', 'Job',
    ↪'Marital', 'Gender', 'Churn',
    ↪'Techie', 'Contract', 'Port_modem', 'Tablet', 'InternetService',
    ↪'Phone', 'Multiple',
    ↪'OnlineSecurity', 'OnlineBackup', 'DeviceProtection',
    ↪'TechSupport', 'StreamingTV',
    ↪
    ↪'StreamingMovies', 'PaperlessBilling', 'PaymentMethod', 'Item1', 'Item2',
    ↪'Item3', 'Item4', 'Item5', 'Item6', 'Item7', 'Item8']

churn_clean_continuous = churn_clean.drop(non_numeric, axis=1)
```

```
[ ]: print(churn_clean_continuous.shape)
```

(10000, 13)

```
[ ]: print(churn_clean_continuous.describe())
```

	Lat	Lng	Population	Children	Age \
count	10000.000000	10000.000000	10000.000000	10000.0000	10000.000000
mean	38.757567	-90.782536	9756.562400	2.0877	53.078400
std	5.437389	15.156142	14432.698671	2.1472	20.698882
min	17.966120	-171.688150	0.000000	0.0000	18.000000
25%	35.341828	-97.082812	738.000000	0.0000	35.000000
50%	39.395800	-87.918800	2910.500000	1.0000	53.000000
75%	42.106908	-80.088745	13168.000000	3.0000	71.000000
max	70.640660	-65.667850	111850.000000	10.0000	89.000000

	Income	Outage_sec_perweek	Email	Contacts \
count	10000.000000	10000.000000	10000.000000	10000.000000
mean	39806.926771	10.001848	12.016000	0.994200
std	28199.916702	2.976019	3.025898	0.988466
min	348.670000	0.099747	1.000000	0.000000
25%	19224.717500	8.018214	10.000000	0.000000
50%	33170.605000	10.018560	12.000000	1.000000
75%	53246.170000	11.969485	14.000000	2.000000
max	258900.700000	21.207230	23.000000	7.000000

	Yearly_equip_failure	Tenure	MonthlyCharge	Bandwidth_GB_Year
count	10000.000000	10000.000000	10000.000000	10000.000000
mean	0.398000	34.526188	172.624816	3392.341550
std	0.635953	26.443063	42.943094	2185.294852
min	0.000000	1.000259	79.978860	155.506715
25%	0.000000	7.917694	139.979239	1236.470827
50%	0.000000	35.430507	167.484700	3279.536903
75%	1.000000	61.479795	200.734725	5586.141370
max	6.000000	71.999280	290.160419	7158.981530

3.2 C2: Standardized Continuous Data Set

```
[ ]: # Import StandardScaler
from sklearn.preprocessing import StandardScaler

scaler = StandardScaler()

# Scaling continuous data set and converting back to data frame using original
↳ column names

scl_churn_clean_continuous = pd.DataFrame(scaler.
↳ fit_transform(churn_clean_continuous), columns= churn_clean_continuous.
↳ columns)

scl_churn_clean_continuous.head()
```

```
[ ]:      Lat      Lng  Population  Children      Age      Income \
0  3.217410 -2.810432  -0.673405 -0.972338  0.720925 -0.398778
1  1.024691  0.431644   0.047772 -0.506592 -1.259957 -0.641954
2  1.213570 -2.142079  -0.417238  0.890646 -0.148730 -1.070885
3 -1.065031 -1.746273   0.284537 -0.506592 -0.245359 -0.740525
4 -1.724710 -0.331512   0.110549 -0.972338  1.445638  0.009478

      Outage_sec_perweek      Email  Contacts  Yearly_equip_failure  Tenure \
0          -0.679978 -0.666282 -1.005852          0.946658 -1.048746
1           0.570331 -0.005288 -1.005852          0.946658 -1.262001
2           0.252347 -0.996779 -1.005852          0.946658 -0.709940
3           1.650506  0.986203  1.017588         -0.625864 -0.659524
4          -0.623156  1.316700  1.017588          0.946658 -1.242551

      MonthlyCharge  Bandwidth_GB_Year
0          -0.003943          -1.138487
1           1.630326          -1.185876
2          -0.295225          -0.612138
3          -1.226521          -0.561857
4          -0.528086          -1.428184
```

```
[ ]: # Exporting scaled DataFrame to CSV
scl_churn_clean_continuous.to_excel('C:/Users/holtb/Data/D212_Data_Mining_II/
↳data/scl_churn_clean_continuous.xlsx')
```

4 Part IV: Analysis

4.1 D1: Matrix of all the principle components

```
[ ]: from sklearn.decomposition import PCA
pca = PCA()

pc = pca.fit_transform(scl_churn_clean_continuous)
pc_df = pd.DataFrame(pc, columns=['PC 1', 'PC 2', 'PC 3', 'PC 4', 'PC 5', 'PC_
↳6', 'PC 7', 'PC 8', 'PC 9', 'PC 10', 'PC 11', 'PC 12', 'PC 13'])

pc_df.head(10)
```

```
[ ]:      PC 1      PC 2      PC 3      PC 4      PC 5      PC 6      PC 7 \
0 -1.632635 -3.296354  0.578184  2.947064  0.049200  0.236283 -1.239584
1 -1.678890 -0.663793 -0.165263 -0.587533 -0.864341  1.227328 -1.145874
2 -0.947985 -1.662107  1.500542  1.260078 -0.604303  1.120935 -1.470020
3 -0.928918  0.883883 -0.255494  1.343020 -1.897874  0.114564  0.424392
4 -1.889784  1.581549 -0.992000  0.972279  1.138263 -0.503190 -0.011660
5 -1.445370  1.578838 -1.434795  0.015113  0.859826  0.043549  0.031879
6 -0.995954  0.066221 -1.083756  0.144645  0.893928  0.764213 -1.899456
7 -1.682875  1.123985  1.616983  0.253361 -0.927824 -1.513094 -0.390062
```

```

8 -1.359716  2.204896  0.950179 -0.243474  2.210810  0.441824 -0.105409
9 -1.827003  0.889171 -0.968896  0.248642  0.220555 -1.578919  0.162032

```

```

      PC 8      PC 9      PC 10      PC 11      PC 12      PC 13
0  0.694516  0.630699  0.126840  0.014701  0.703792 -0.025562
1  1.079940 -0.122828  0.062935 -1.462859  0.927392 -0.036908
2 -0.081259  0.009108  0.484669  0.729901 -0.259270  0.060309
3 -1.721621  0.296289 -0.700928  0.262978 -1.253845  0.127564
4 -0.813185  1.456770  0.191183 -0.479151 -1.201817 -0.058608
5 -1.073229  1.230639  1.684685  0.274779 -0.277542 -0.019170
6  0.601758 -0.125758 -0.599199 -0.035029 -0.487417  0.108837
7 -0.442926  0.166403 -0.470777 -0.616362  0.081043  0.064701
8 -2.030327  3.203556  0.913609 -2.648496 -0.460559  0.081075
9  0.149250  1.662681 -0.503430  0.393250  0.637253 -0.045272

```

4.2 D2: Total Number of Principle Components Using Elbow Rule

```

[ ]: print(len(pca.components_))

var = pca.explained_variance_ratio_

```

13

```

[ ]: import matplotlib.pyplot as plt


plt.plot(var)
plt.xlabel('Principal component index')
plt.ylabel('Explained variance ratio')

```

```

[ ]: Text(0, 0.5, 'Explained variance ratio')

```



```
[ ]: # PCA fit with the last 2 principle components removed
pca =PCA(n_components=11)

pc = pca.fit_transform(scl_churn_clean_continuous)
pc_df = pd.DataFrame(pc, columns=['PC_1', 'PC_2', 'PC_3', 'PC_4', 'PC_5', 'PC_6', 'PC_7', 'PC_8', 'PC_9', 'PC_10', 'PC_11'])

pc_df.head()
```

```
[ ]:      PC_1      PC_2      PC_3      PC_4      PC_5      PC_6      PC_7 \
0 -1.632635 -3.296354  0.578184  2.947064  0.049200  0.236283 -1.239584
1 -1.678890 -0.663793 -0.165263 -0.587533 -0.864341  1.227328 -1.145874
2 -0.947985 -1.662107  1.500542  1.260078 -0.604303  1.120935 -1.470020
3 -0.928918  0.883883 -0.255494  1.343020 -1.897874  0.114564  0.424392
4 -1.889784  1.581549 -0.992000  0.972279  1.138263 -0.503190 -0.011660

      PC_8      PC_9      PC_10      PC_11
0  0.694516  0.630699  0.126840  0.014701
1  1.079940 -0.122828  0.062935 -1.462859
2 -0.081259  0.009108  0.484669  0.729901
3 -1.721621  0.296289 -0.700928  0.262978
4 -0.813185  1.456770  0.191183 -0.479151
```

4.3 D3: Variance of Each Principle Component

```
[ ]: print(pca.explained_variance_)

[1.99490495 1.23415146 1.05376978 1.04469451 1.02328808 1.01078182
 0.99930258 0.99267518 0.98453106 0.9618864 0.96049531]

[ ]: print(pca.explained_variance_ratio_)

[0.15343888 0.09492523 0.08105111 0.08035308 0.0787066 0.07774467
 0.07686174 0.07635199 0.07572558 0.07398386 0.07387687]
```

4.4 D4: Total Variance

```
[ ]: print(pca.explained_variance_.cumsum())

[ 1.99490495 3.22905642 4.28282619 5.3275207 6.35080878 7.3615906
 8.36089318 9.35356836 10.33809941 11.29998582 12.26048113]

[ ]: print(pca.explained_variance_ratio_.cumsum())

[0.15343888 0.24836412 0.32941522 0.4097683 0.4884749 0.56621957
 0.64308131 0.71943331 0.79515889 0.86914276 0.94301962]
```

4.5 D5: Summarize Results

Using `sklearn`'s PCA method the standardized data set was fit and transformed. The `explained_variance_ratio_` feature of the PCA was plotted against the indexed principle components (PC) to create the scree plot. After analyzing the scree plot it was difficult to determine an elbow in the plot. This was due to the fact that most of the PCs explained around 7-9% variance for each one. The exceptions were PC1, explaining ~15% of variance and PC 13 explaining almost 0% of the variance. The final PC was conducted with only 2 PC's removed, however, during model fitting, further experimentation with PC removal is needed to produce the most efficient model.

Each PC can be explored by looking at the effects of each variable on the PC from the table created before. For simplicity the first and second PC are explored. For PC_1 it can be determined that 'Tenure' and 'Bandwidth_GB_Year' have the greatest effect in the principle component as their relationship is very strong. Having studied the data previously it is known that both of these variables have a strong linear relationship. For the second principle component, 'Population', 'Lng' and 'Email' have a moderate relationship in the data. This is most likely caused because the geographic locations in the data set are skewed because many of the data points are located much further 'east' (continental US) compared to another portion of the dataset. (Alaska & Hawaii)

```
[ ]: vectors = pca.components_.round(2)

[ ]: # Create dictionary of column names to PC components
PC_1_effects = dict(zip(scl_churn_clean_continuous.columns, vectors[0]))
PC_2_effects = dict(zip(scl_churn_clean_continuous.columns, vectors[1]))
PC_3_effects = dict(zip(scl_churn_clean_continuous.columns, vectors[2]))
PC_4_effects = dict(zip(scl_churn_clean_continuous.columns, vectors[3]))
PC_5_effects = dict(zip(scl_churn_clean_continuous.columns, vectors[4]))
```

```

PC_6_effects = dict(zip(scl_churn_clean_continuous.columns, vectors[5]))
PC_7_effects = dict(zip(scl_churn_clean_continuous.columns, vectors[6]))
PC_8_effects = dict(zip(scl_churn_clean_continuous.columns, vectors[7]))
PC_9_effects = dict(zip(scl_churn_clean_continuous.columns, vectors[8]))
PC_10_effects = dict(zip(scl_churn_clean_continuous.columns, vectors[9]))
PC_11_effects = dict(zip(scl_churn_clean_continuous.columns, vectors[10]))

```

```

[ ]: # Create summary data frame of PC effects
PC_effects = pd.DataFrame([PC_1_effects, PC_2_effects, PC_3_effects,
    ↪PC_4_effects, PC_5_effects, PC_6_effects, PC_7_effects, PC_8_effects,
    ↪PC_9_effects, PC_10_effects, PC_11_effects], index= ['PC_1',
    ↪'PC_2', 'PC_3', 'PC_4', 'PC_5', 'PC_6', 'PC_7', 'PC_8', 'PC_9', 'PC_10',
    ↪'PC_11'])
PC_effects

```

```

[ ]:
      Lat  Lng  Population  Children  Age  Income  Outage_sec_perweek  \
PC_1 -0.02  0.01      -0.00      0.01  0.00      0.00      0.01
PC_2 -0.71  0.18       0.65     -0.01  0.05     -0.05      0.01
PC_3 -0.03 -0.29       0.15      0.45 -0.44      0.20     -0.25
PC_4  0.11 -0.74       0.32     -0.46  0.23     -0.04     -0.13
PC_5 -0.09  0.34      -0.12     -0.11  0.44      0.31     -0.46
PC_6 -0.03 -0.09       0.10      0.13 -0.10      0.10      0.60
PC_7 -0.01 -0.05       0.05      0.03 -0.19      0.77      0.05
PC_8 -0.02 -0.09       0.08     -0.07  0.09      0.34     -0.18
PC_9  0.09 -0.17      -0.03      0.19  0.34      0.25      0.06
PC_10 0.02 -0.15       0.06      0.44 -0.08     -0.28     -0.52
PC_11 0.05 -0.11       0.10      0.57  0.61     -0.03      0.22

```

```

      Email  Contacts  Yearly equip failure  Tenure  MonthlyCharge  \
PC_1 -0.02      0.00      0.02      0.71      0.04
PC_2  0.15      0.03     -0.01     -0.01     -0.00
PC_3 -0.09     -0.45      0.15      0.01     -0.40
PC_4 -0.14      0.11      0.06      0.03     -0.14
PC_5 -0.35      0.01      0.42      0.01     -0.22
PC_6 -0.40      0.08      0.59     -0.04      0.26
PC_7  0.00      0.52     -0.29     -0.00     -0.04
PC_8 -0.13     -0.51     -0.19     -0.04      0.71
PC_9  0.76     -0.05      0.40      0.00      0.06
PC_10 -0.05      0.49      0.14     -0.04      0.41
PC_11 -0.25     -0.03     -0.38      0.01     -0.14

```

```

      Bandwidth_GB_Year
PC_1      0.71
PC_2     -0.01
PC_3      0.01
PC_4     -0.00
PC_5     -0.02

```


PC_6	-0.01
PC_7	0.00
PC_8	0.00
PC_9	0.00
PC_10	0.01
PC_11	-0.01

5 Part V: Attachments

5.1 Code Sources

Boeye, J. “Dimensionality Reduction in Python” [MOOC]. Datacamp. <https://app.datacamp.com/learn/courses/dimensionality-reduction-in-python>

Harris, C.R., Millman, K.J., van der Walt, S.J. et al. Array programming with NumPy. *Nature* 585, 357–362 (2020). DOI: 0.1038/s41586-020-2649-2. (Publisher link).

J. D. Hunter, “Matplotlib: A 2D Graphics Environment”, *Computing in Science & Engineering*, vol. 9, no. 3, pp. 90-95, 2007.

Pedregosa, F., Varoquaux, G., Gramfort, A., Michel, V., Thirion, B., Grisel, O., Blondel, M., Prettenhofer, P., Weiss, R., Dubourg, V., Vanderplas, J., Passos, D., Brucher, M., Perrot, M., & Duchesnay, E. (2011). Scikit-learn: Machine Learning in Python. *Journal of Machine Learning Research*, 12, 2825–2830.

Python Software Foundation. Python Language Reference, version 3.7. Available at <http://www.python.org>

5.2 In-text Citations

Laerd Statistics. (2018). Principal Components Analysis (PCA) using SPSS Statistics. Retrieved January 02, 2022, from <https://statistics.laerd.com/spss-tutorials/principal-components-analysis-pca-using-spss-statistics.php>